

Configuring Connectors

Connector configurations are key-value mappings. For standalone mode these are defined in a properties file and passed to the Connect process on the command line. In distributed mode, they will be included in the JSON payload sent over the [REST API](#) for the request that creates (or modifies) the connector. When choosing to deploy in standalone or distributed mode, it is useful to review the recommendations [here](#). Both standalone and distributed connectors can be configured by using the [Confluent Control Center](#) which gives a graphical interface to update connector configurations. Most configurations are connector dependent, but there are a few settings common to all connectors:

- `name` - Unique name for the connector. Attempting to register again with the same name will fail.
- `connector.class` - The Java class for the connector
- `tasks.max` - The maximum number of tasks that should be created for this connector. The connector may create fewer tasks if it cannot achieve this level of parallelism.
- `key.converter` - (optional) Override the default key converter class set by the worker.
- `value.converter` - (optional) Override the default value converter class set by the worker.

To pass configuration parameters to key and value converters, prefix them with `key.converter.` or `value.converter.` as you would in the worker configuration when defining default converters. Note that these are only used when the corresponding converter configuration is specified in the `key.converter` or `value.converter` properties.

Note that these parameters are not used unless the corresponding converter configuration is specified in the connector configuration.

Sink connectors also have one additional option to control their input:

- `topics` - A list of topics to use as input for this connector

For other options, consult the documentation for [individual connectors](#).

Standalone Example

Below is an example of a standalone connector configuration for the supported `FileStreamSinkConnector`. Note the common configuration `name`, `connector.class`, `tasks.max`, and `topics`, and one `FileStreamSinkConnector` specific configuration `file` is specified. The file containing this configuration should be added as shown [here](#).

```
name=local-file-sink
connector.class=FileStreamSinkConnector
tasks.max=1
file=test.sink.txt
topics=connect-test
```

Distributed Example

You can use the REST API to manage the connectors running on workers in distributed mode. Here is a simple example that creates a `FileStreamSinkConnector` as in the standalone example. Note the URL is pointing to localhost indicating that the connect worker has been started on the same host where the `curl` command is run. Instructions for starting a worker in distributed mode are [here](#).

```
curl -X POST -H "Content-Type: application/json" --data '{"name": "local-file-sink", "config": {"connector.class": "FileStreamSinkConnector", "tasks.max": "1", "file": "test.sink.txt", "topics": "connect-test" }}' http://localhost:8083/connectors
# Or, to use a file containing the JSON-formatted configuration
# curl -X POST -H "Content-Type: application/json" --data @config.json http://localhost:8083/connectors
```

Expand
Content

To create a connector, you start the workers and then make a REST request to create a connector as above. Unlike many other systems, all nodes in Kafka Connect can respond to REST requests, including creating, listing, modifying, and destroying connectors (see the [REST API](#) section for details).

© Copyright 2019, Confluent, Inc. [Privacy Policy](#) | [Terms & Conditions](#). Apache, Apache Kafka, Kafka and the Kafka logo are trademarks of the [Apache Software Foundation](#). All other trademarks, servicemarks, and copyrights are the property of their respective owners.

[Please report any inaccuracies on this page or suggest an edit.](#)

21 Votes



Last updated on Oct 17, 2019.