

# Application Reset Tool

You can reset an application and force it to reprocess its data from scratch by using the application reset tool. This can be useful for development and testing, or when fixing bugs.

The application reset tool handles the Kafka Streams [user topics](#) (input, output, and intermediate topics) and [internal topics](#) differently when resetting the application.

Here's what the application reset tool does for each topic type:

- Input topics: Reset offsets to specified position. By default they are reset to the beginning of the topic.
- Intermediate topics: Skip to the end of the topic, i.e., set the application's committed consumer offsets for all partitions to each partition's `logSize` (for consumer group `application.id`).
- Internal topics: Delete the internal topic (this automatically deletes any committed offsets).

The application reset tool does not:

- Reset output topics of an application. If any output (or intermediate) topics are consumed by downstream applications, it is your responsibility to adjust those downstream applications as appropriate when you reset the upstream application.
- Reset the local environment of your application instances. It is your responsibility to delete the local state on any machine on which an application instance was run. See the instructions in section [Step 2: Reset the local environments of your application instances](#) on how to do this.

## Prerequisites

- All instances of your application must be stopped. Otherwise, the application may enter an invalid state, crash, or produce incorrect results. You can verify whether the consumer group with ID `application.id` is still active by using `bin/kafka-consumer-groups`.
- Use this tool with care and double-check its parameters: If you provide wrong parameter values (e.g., typos in `application.id`) or specify parameters inconsistently (e.g., specify the wrong input topics for the application), this tool might invalidate the application's state or even impact other applications, consumer groups, or your Apache Kafka® topics.
- You should manually delete and re-create any intermediate topics before running the application reset tool. This will free up disk space in Kafka brokers.
- You should delete and recreate intermediate topics before running the application reset tool, unless the following applies:
  - You have external downstream consumers for the application's intermediate topics.
  - You are in a development environment where manually deleting and re-creating intermediate topics is unnecessary.

## Step 1: Run the application reset tool

Invoke the application reset tool from the command line

```
<path-to-confluent>/bin/kafka-streams-application-reset
```

The tool accepts the following parameters:

Option (* = required)	Description
* --application-id <String: ID>	The  kstreams  application ID (application.id).
--bootstrap-servers <String: urls>	Comma-separated list of broker urls with format: HOST1:PORT1,HOST2:PORT2 (default: localhost:9092)
--by-duration <String: urls>	Reset offsets to offset by duration from current timestamp. Format: 'PnDTnHnMnS'
--config-file <String: file name>	Property file containing configs to be passed to admin clients and embedded consumer.
--dry-run	Display the actions that would be performed without executing the reset commands.
--from-file <String: urls>	Reset offsets to values defined in CSV file.
--input-topics <String: list>	Comma-separated list of user input topics. For these topics, the tool will reset the offset to the earliest available offset.
--intermediate-topics <String: list>	Comma-separated list of intermediate user topics (topics used in the through() method). For these topics, the tool will skip to the end.
--shift-by <Long: number-of-offsets>	Reset offsets shifting current offset by 'n', where 'n' can be positive or negative
--to-datetime <String>	Reset offsets to offset from datetime. Format: 'YYYY-MM-DDTHH:mm:ss.sss'
--to-earliest	Reset offsets to earliest offset.
--to-latest	Reset offsets to latest offset.
--to-offset <Long>	Reset offsets to a specific offset.
--zookeeper	Zookeeper option is deprecated by bootstrap.servers, as the reset tool would no longer access Zookeeper directly.

Parameters can be combined as needed. For example, if you want to restart an application from an empty internal state, but not reprocess previous data, simply omit the parameters `--input-topics` and `--intermediate-topics`.

## Step 2: Reset the local environments of your application instances

For a complete application reset, you must delete the application's local state directory on any machines where the application instance was run. You must do this before restarting an application instance on the same machine. You can use either of these methods:

- The API method `KafkaStreams#cleanup()` in your application code.
- Manually delete the corresponding local state directory (default location: `/var/lib/kafka-streams/<application.id>`). For more information, see [state.dir](#) StreamsConfig class.

## Example

In this example you are developing and testing an application locally and you want to iteratively improve your application via run-reset-modify cycles.

```

package io.confluent.examples.streams;

import ...;

public class ResetDemo {

    public static void main(String[] args) throws Exception {
        // Kafka Streams configuration
        Properties props = new Properties();
        props.put(StreamsConfig.APPLICATION_ID_CONFIG, "my-streams-app");
        // ...and so on...

        // Define the processing topology
        StreamsBuilder builder = new StreamsBuilder();
        builder.stream("my-input-topic")
            .selectKey(...)
            .through("rekeyed-topic")
            .groupByKey()
            .count("global-count")
            .to("my-output-topic");

        KafkaStreams app = new KafkaStreams(builder.build(), props);

        // Delete the application's local state.
        // Note: In real application you'd call `cleanup()` only under
        // certain conditions. See tip on `cleanup()` below.
        app.cleanup();

        app.start();

        // Note: In real applications you would register a shutdown hook
        // that would trigger the call to `app.close()` rather than
        // using the sleep-then-close example we show here.
        Thread.sleep(30 * 1000L);
        app.close();
    }
}

```

**Tip:** To avoid the corresponding recovery overhead, you should not call `cleanup()` unconditionally and every time an application instance is restarted or resumed. For example, in a production application you could use command line arguments to enable or disable the `cleanup()` call on an as-needed basis.

You can then perform run-reset-modify cycles as follows:

```

# Run your application
bin/kafka-run-class io.confluent.examples.streams.ResetDemo

# After stopping all application instances, reset the application
bin/kafka-streams-application-reset --application-id my-streams-app \
    --input-topics my-input-topic \
    --intermediate-topics rekeyed-topic

# Now you can modify/recompile as needed and then re-run the application again.
# You can also experiment, for example, with different input data without
# modifying the application.

```

---

© Copyright 2019, Confluent, Inc. [Privacy Policy](#) | [Terms & Conditions](#). Apache, Apache Kafka, Kafka and the Kafka logo are trademarks of the [Apache Software Foundation](#). All other trademarks, servicemarks, and copyrights are the property of their respective owners.

Please report any inaccuracies on this page or suggest an edit.

Rate this page



Last updated on Sep 10, 2019.